



Opera Solutions, LLC
180 Maiden Lane
17th Floor
New York, NY 10038
+1 (646) 437 2100 telephone
+1 (646) 437 2101 facsimile
www.operasolutions.com

Personalize Expedia Hotel Searches – ICDM 2013

High-Level approach

Dec 6th, 2013

NOTICE: Proprietary and Confidential

This material is proprietary to Opera Solutions. It contains trade secrets and confidential information which is sole property of Opera Solutions. This material is solely for the Client's internal use. This material shall not be used, reproduced, copied, disclosed, transmitted, in whole or in part, without the express written consent of Opera Solutions.

© 2013 Opera Solutions, LLC. All rights reserved.



Individual models

Performance check (model validation)

- 10% random subset as validation set fixed for ALL experiments

SGD model

- Own implementation
- Linear model on 19 numeric features (total model has only 19 parameters !)
- Trained with pairwise ranking updates for 50 epochs
- 0.50377 on leaderboard

Neural Network

- Own implementation
- Inputs: numeric + categoric features
- 1 hidden layer 30 neurons (tanh)
- Minibatch training with rank-updates
- Best individual: 0.5297 on leaderboard

GB - Gradient Boosted Decision Tree For Ranking

- Own implementation
- Base idea from C. Burges's LambdaMART
- Added: support for categoric features, random splits
- Best individual: 0.5256 on leaderboard

LambdaMART by RankLib

- → <http://sourceforge.net/p/lemur/wiki/RankLib/>
- Inputs: numeric features + average value of those per prop_id
- Best individual: 0.5338 on leaderboard

Model	NDCG valid	NDCG leaderboard
Random benchmark	0.348	0.3479
Position benchmark	0.4967	0.4999
valueBook sum	0.443	0.4418
SVD on prop_id / srch_dest_id	0.4618	0.46317
Good consistency of validation and leaderboard Score by taking a random 10% data subset		

Blending

24 individual predictors (leaderboard scores)

- 6x Neural Net (0.511 ... 0.529)
- 1x SGD model (0.52)
- 4x LambdaMART (0.510 ... 0.533)
- 11x GB - Gradient Boosted Decision Tree For Ranking (0.49 ... 0.526)

Linear Blender
$$s = \sum_{i=1}^N p_i w_i$$

- Weights are search by parameter searcher APT1 (see Netflix Prize doc)
- Maximizing the NDCG on the validation set
- Take the best runs from 100 random initializations
 - Problems with many local minima
- Final solution: 0.5407 leaderboard score

Insights / Conclusion

- An ensemble of ranking models outperforms the best individual
- Dataset was large enough to use a single validation set (no cross-validation needed)
- Good data quality
- Time and price compare features does not help
- Biggest improvements by adding features for each prop_id
 - Mean of numeric values
 - Stddev of numeric values
 - Median of numeric values

